Exploring the Evolution and Applications of Geospatial Databases: A Comprehensive Review

1st Mohsen Jafari *Computer Science*
Bowling Green State University
Bowling Green, Ohio, USA
mjafari@bgsu.edu

***Abstract*—Geospatial databases have become increasingly vital in managing and analyzing spatially-referenced data, with applications across various domains. This paper consolidates insights from multiple abstracts, offering a comprehensive overview of the evolution, capabilities, and challenges of geospatial databases. It discusses the emergence of geospatial databases from traditional paper maps to digital formats, exploring relational, spatial, and NoSQL databases and their suitability for different scenarios. The review also addresses key challenges such as data interoperability and scalability, emphasizing the need for efficient solutions for managing geospatial big data. Additionally, it examines emerging trends like geospatial video applications and the integration of geospatial data in online analysis tools. By synthesizing these diverse perspectives, this abstract contributes to understanding the potential of geospatial databases in addressing complex spatial challenges and driving innovation across industries.**

***Keywords*—Geospatial Database, Big Data, NoSQL, Geospatial Analytics, Store Data.**

I. Introduction

Geospatial databases are specialized systems designed to store, manage, and analyze spatially-referenced information, such as maps, satellite imagery, and location-based data. They provide a structured way to organize geographical data, allowing users to query and retrieve information based on its spatial characteristics. These databases play a crucial role in various fields, including urban planning, environmental management, and geographic information systems (GIS). Unlike traditional databases, geospatial databases have features tailored to handle spatial data efficiently, such as spatial indexing techniques and specialized query optimization strategies. They can store different types of spatial data, including points, lines, and polygons, along with associated attributes. Geospatial databases enable users to perform complex spatial analyses, such as proximity analysis, spatial interpolation, and spatial modeling. They are essential tools for decision-making processes in areas such as disaster management, transportation planning, and natural resource management. Geospatial databases also support the integration of spatial data from multiple sources, facilitating interoperability and data sharing among different organizations and systems.

In recent years, the proliferation of geospatial applications has led to an exponential increase in the variety, volume, and velocity of spatial data. From satellite imagery to user-generated content, the sheer abundance of geolocated data presents both unprecedented opportunities and daunting challenges. Geospatial databases, serving as the backbone of modern data management in spatially-referenced information, have emerged as essential tools for storing, querying, and analyzing this vast and diverse array of data.

This paper aims to provide a comprehensive overview of geospatial databases, exploring their evolution, capabilities, and practical applications across diverse domains. As the world transitions from traditional paper maps to digital formats, the role of geospatial databases has evolved significantly, offering sophisticated solutions for managing complex spatial data. From Geographic Information Systems (GIS) to urban planning and environmental management, the versatility and utility of geospatial databases permeate numerous fields, driving innovation and addressing complex spatial challenges.

The purpose of this paper id delving into the core components of geospatial databases, examining their architectural frameworks, data models, and standards. The discussion encompasses various types of geospatial databases, including relational, spatial, and NoSQL solutions, elucidating their features, advantages, and limitations. Additionally, we explore the challenges inherent in managing geospatial big data, such as scalability and data interoperability, and investigate emerging trends and future directions in geospatial database research.

Geospatial data, which refers to data about the surface of the Earth, has become increasingly important in recent years due to advancements in technologies like satellites, drones, GPS, and location-based services. The huge volumes of geospatial big data being generated, collected and used present major challenges for data management and analysis systems. Traditionally, relational database management systems (RDBMS) have been used to store and process geospatial data, but they face limitations when dealing with geospatial big data that has high volume, high velocity of changes, and a variety of data structures.

The characteristic volume, variety, velocity and other dimensions like variability, veracity and visualization add complexities in storing, managing, processing and serving geospatial big data. Ensuring secure access, defining user access policies, supporting spatial queries like range queries and nearest neighbor searches are some of the key requirements for an effective geospatial database system. As the amount of geospatial data continues growing exponentially, there is a crucial need for efficient systems capable of handling and extracting insights from these large geospatial datasets to maximize understanding of associated phenomena and facilitate data-driven decision making.

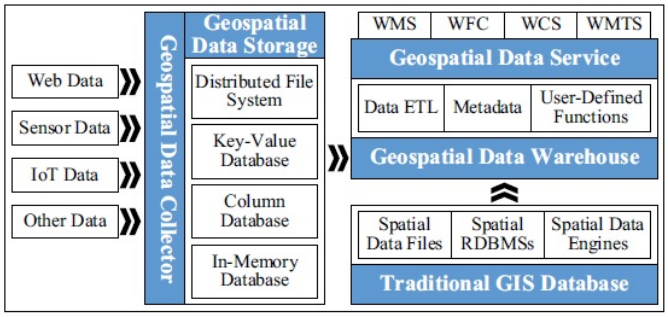

Fig. 1. An overview of the GeoSpatial Data Store

To address these challenges, solutions like NoSQL databases and distributed processing frameworks like Hadoop and Spark have emerged as alternatives or complements to traditional RDBMS. These modern approaches provide capabilities to handle the scale, speed, and flexibility required for geospatial big data applications like serving standard geospatial web services efficiently. Developing comprehensive data models, spatial indexing techniques, security policies and scalable architectures are important research areas for building high-performance geospatial database systems.

The Geospatial Data Store, depicted in Figure 1, aggregates data from diverse sources like Sensor Web, IoT, or the Web. It proposes a geospatial data warehouse to manage heterogeneous datasets, employing advanced storage paradigms. Leveraging HBase, scalable big data storage on HDFS, it facilitates standardized services for geospatial big data access. Elasticsearch creates secondary indexes on HBase tables for efficient spatial, temporal, attribute, and multi-conditional queries. Integration of Apache Impala enables BI/analytic queries, simplifying the extension of GeoServer for standardized geospatial data services.

II. Related Work

The growth of geospatial big data has necessitated robust solutions for its management, processing, and secure access control. The paper [7]. provides an overview of geospatial big data characteristics like volume, variety and velocity, and evaluates the performance of open source tools like SpatialHadoop and GeoSpark for processing such data, concluding that GeoSpark outperforms on most metrics. Complementing this technological perspective, the paper [8] examines the security concerns surrounding sensitive geospatial data. It reviews spatial data modeling methods, existing access control policies like role-based and fine-grained, and proposes a comprehensive data model incorporating robust security policies for ensuring confidentiality, integrity and controlled access to geospatial data.The paper [11] proposes a geospatial big data platform for providing geospatial information services. It discusses the key components of the platform, including data storage, data management, spatial analysis, and application services. The platform aims to handle large volumes of geospatial data and provide efficient spatial query processing, data visualization, and application development. The proposed architecture includes components such as GeoSpark, GeoTrellis, GeoMesa, and GeoWave for distributed spatial data processing and analysis.

This paper [5] presents Apperception, a new type of database management system designed for geospatial video data. Apperception provides an easy-to-use data model and programming interface that allows developers to collectively reason about entities observed across multiple geospatial video streams. The paper demonstrates Apperception through a web interface for an AMBER Alert application for vehicle tracking. Users can ingest multiple video datasets, select a vehicle of interest (and optionally specify spatiotemporal constraints), and Apperception will overlay the vehicle's trajectory on a map and show cropped video of the selected object. The key technical contribution highlighted is the query language that allows querying objects across a multi-camera world.

The paper [12] proposes the Geo-Planar Indexing (GPI) scheme for efficient retrieval of raster-based geospatial data from databases for mobile GIS applications. GPI assigns unique 8-digit index numbers to small tiles partitioning a base map image. Compared to R-Tree indexing, GPI provides faster retrieval performance of 1.3 to 4.3 times across different data densities. GPI operates under standard RDBMS without requiring specialized spatial extensions. It enables rapid delivery of geo-referenced map data to moving clients in mobile GIS scenarios.

The article [11] paper discusses the Open Geospatial Datastore Interface (OGDI), which is a public-domain application programming interface (API) that allows software applications to directly read and access various geospatial data formats without translation. OGDI aims to address the geospatial data barrier by providing a multilingual approach to reading different data formats. The paper outlines OGDI's architecture, transient data model, available drivers, and how to use the API within applications. OGDI facilitates direct and transparent access to geospatial data, enabling interoperability among different geospatial databases and formats. The Figure 2 Shows the architecture of the OGDI.

The paper [13] discusses the concept of geospatial knowledge discovery (KDG), which is a branch of knowledge discovery from spatial databases. It presents a framework for spatial knowledge management (SKM) and spatial knowledge management systems (SKMS) to facilitate KDG. The paper also provides a case study on designing a precision agriculture knowledge management and discovery system based on geographic information systems (GIS), global positioning systems (GPS), and remote sensing (RS) technologies, collectively known as "3S" integration. The Figure 3 illustrates the general process of spatial knowledge management.

The paper [2] compares Azure SQL Database and Azure DocumentDB for geospatial applications, assessing performance, scalability, and cost-effectiveness. It evaluates response time, concurrent user impact, and record volume on database performance, along with functionality like geometry object support and spatial index support. Results suggest DocumentDB is generally faster, while SQL Database scales better with concurrent users. It aims to aid in selecting the best-suited database for geospatial needs and acknowledges funding from the European Union's Horizon 2020 program.

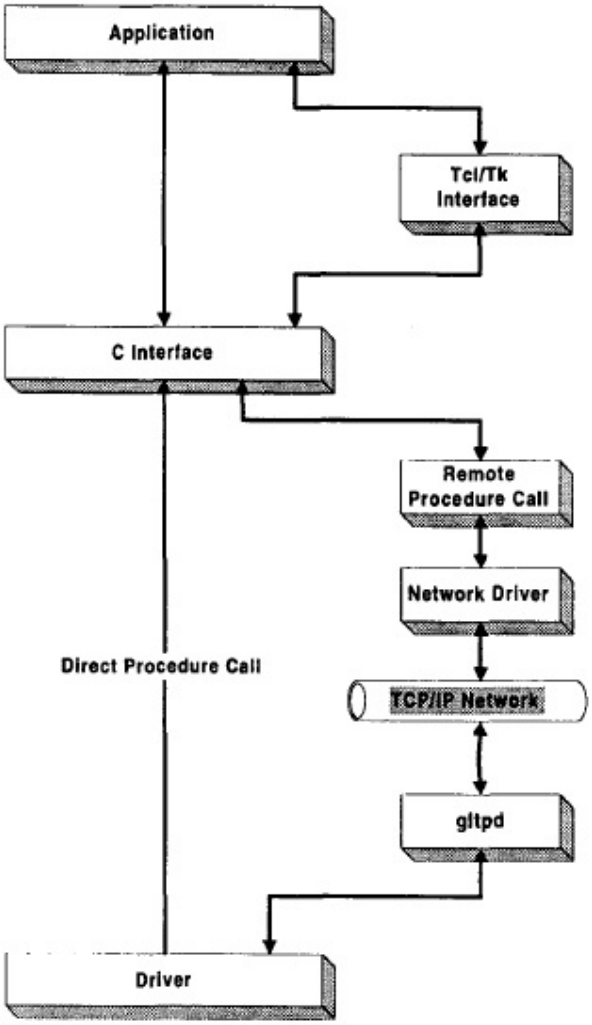

Fig. 2. OGDI Architecture

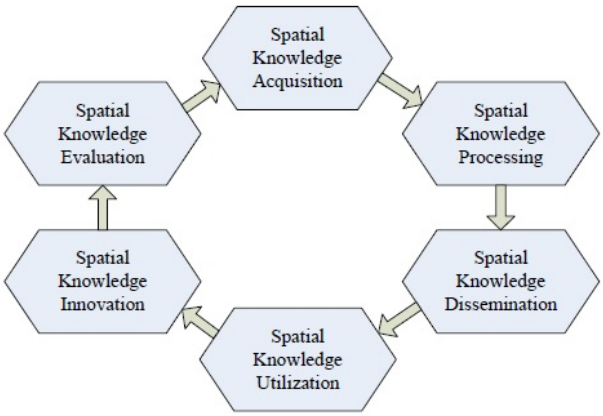

Fig. 3. General process of spatial knowledge management

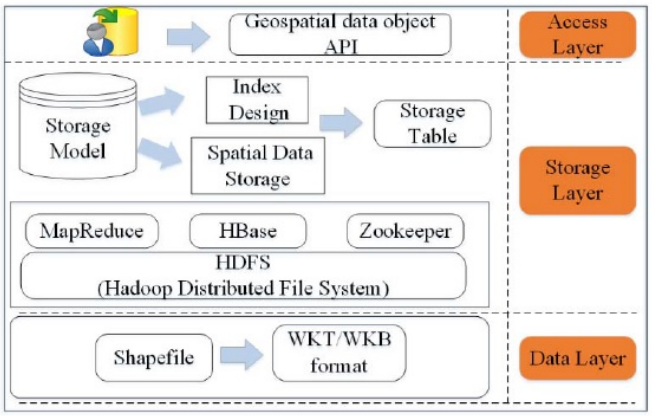

Fig. 4. The framework of geospatial data storage

The article [9] delves into in-database geospatial analytics challenges and solutions with Python, emphasizing the surge in spatial data and the necessity for efficient analysis tools. It introduces the ibmdbpy-spatial framework, enabling spatial analysis within IBM DB2, tackling complex spatial queries and data heterogeneity. The framework facilitates fast spatial analysis without data storage in-memory, demonstrated through a case study on crime hotspots in New York City. Overall, it provides a scalable, user-friendly approach for in-database analytics, with implications for smart city development and decision-making processes.

The article [14] surveys measures and methods for matching geospatial datasets, providing a simple classification, two taxonomies and practical issues for choosing measures and methods.

The paper [4] discusses the limitations of traditional relational database management systems (RDBMS) in storing and analyzing big data, particularly geospatial data. The authors propose a new method for geospatial data storage using HBase and MapReduce, which improves I/O efficiency and manages data effectively. They demonstrate the efficiency of their method by importing a large dataset into HBase using MapReduce, which significantly reduces the time required compared to traditional methods. The Figure 4 shows the process of storing geospatial data.

The paper [3] discusses the development of an online platform for managing geospatial resources. The platform, named Open Geospatial Engine (OGE), integrates two core components: the Data Center for data integration and storage, and the Model Center for registration and sharing of geospatial analysis models. The platform leverages cloud service technologies and adheres to Open Geospatial Consortium (OGC) standard API interfaces. The system's applicability is demonstrated through a case study involving the discovery, access, and manipulation of large datasets and models.

The paper [10] addresses challenges in geospatial data processing posed by the proliferation of high-resolution remote sensing data. It presents a novel framework utilizing GPU clusters to enhance processing throughput, employing both CPU and GPU parallelism to concurrently process sub-rasters. The development of a multilevel parallel geospatial data processing

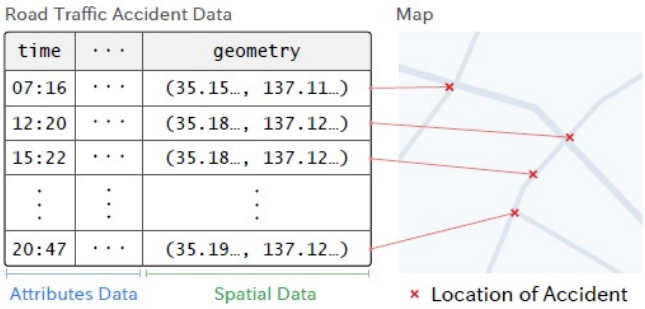

Fig. 5. Example of traffic accidents as geospatial data.

framework is highlighted, showcasing its application in EO raster processing and DEM information mining, resulting in significant speed-ups compared to CPU-only methods.

The paper [1] discusses the creation of a comprehensive database for analyzing road traffic accidents. The authors highlight the need for a database that can handle spatiotemporal data, given that traffic accidents are spatiotemporal events. The designed database allows for the addition of various attributes and links to different data, enabling a multifaceted approach to traffic accident analysis. The database was implemented using PostgreSQL and PostGIS, and its performance was evaluated using real-world traffic accident data. The results demonstrated that the database has good operational performance and is useful for traffic accident data analysis, thus meeting the requirements for a data model. The authors suggest that the database could be further improved by considering spatial indexes and by continually examining relevant data. The Figure 5 illustrates the traffic accidents as geospatial data.

The paper [6] presents GeoYCSB, a novel benchmark framework designed for evaluating the performance and scalability of geospatial NoSQL databases. The authors use this framework to assess the capabilities of MongoDB and Couchbase, focusing on the influence of various factors such as the number of concurrent users, the number of records, and the configuration cost on the average response time of queries. The paper also includes a comparative analysis of the efficiency of ArsGIS and MongoDB in retrieving a specified number of location points within a fixed bounding box. Furthermore, the authors introduce SpatialIgnite, an extension of the Apache Ignite data system that incorporates geospatial operations, and conduct benchmarking experiments to compare its performance with SpatialHadoop and GeoSpark. The paper fills a gap in the literature by providing a benchmarking tool for replicated and sharded NoSQL databases for geospatial workloads.

III. Geospatial Database Usage and methodology

In this paper, our aim is to delve deeply into the intricate world of geospatial databases, exploring their functionality, applications, and the myriad ways they shape our understanding of spatial data. We will meticulously examine the utilization of geospatial databases across various domains, shedding light on their indispensable role in modern data management. Our investigation will encompass a comprehensive analysis of pertinent literature, spanning topics such as the integration of geospatial data within NoSQL databases, efficient methods for storing and retrieving spatial information, programming paradigms tailored for geospatial data processing, advanced indexing techniques to accelerate spatial queries, the synergy between geospatial databases and big data analytics, as well as the evolving landscape of visualization tools tailored for spatial data exploration. Furthermore, we will explore the emerging trend of leveraging GPU (Graphics Processing Unit) technology to enhance the processing speed and efficiency of geospatial data operations, illustrating how these advancements are reshaping the frontier of geospatial database management. Through this exploration, we aim to provide a comprehensive understanding of the multifaceted realm of geospatial databases and their profound impact on modern data ecosystems.

A. *Definition and characteristics of geospatial big data*

The role of big data and geographic information systems (GIS) in addressing global problems such as climate change and pandemics are identified. The integration of geospatial big data is crucial for consolidating datasets related to healthcare, the environment, and other domains. The big data integration processes classifies into three categories: data warehousing, data transformation, and integration methods. It also highlights research challenges, such as the need for effective data integration, semantic interoperability, and the processing and storage of geospatial data.The Figure 6 shows the classification of big data integration with geospatial data.

1) Data Warehousing: The evolution of data warehousing (DW) from traditional servers to modern architectures, including the classification of classical and modern data warehouses. The challenges of handling big data in traditional DWs and the emergence of modern DWs designed to handle large volumes of structured, semi-structured, and unstructured data. Additionally, it covers spatial data warehouses (SDWs) and data cubes as effective approaches for storing and analyzing big earth data. The data warehousing delves into the ETL (extract, transform, load) process, which is crucial in the creation of DWs, and the need for innovative techniques to integrate heterogeneous geospatial and non-geospatial data. It is important to integrate GIS technologies with data warehousing to obtain SDWs and the development of spatial ETL processes to handle massive amounts of geospatial data.

2) Data Transformation: The data transformation discusses the Extract-Transform-Load (ETL) process, which is crucial for creating a data warehouse (DW) and integrating data from different sources. Traditional ETL is an eager process that loads complete data sources into an integrated repository, while lazy ETL involves loading only metadata. The data transformation also highlights the limitations of traditional ETL in handling large datasets and the need for new approaches, such as cloud computing, MapReduce, and NoSQL data models. Also, it addresses the challenges of integrating real-time data with archive data and the need for distributed ETL processes to promote integrated data sharing. The data transformation emphasizes the importance of re-examining

ETL to address the complex nature of big data and the need for new approaches to optimize its efficacy.

3) Integration Methods: The role of domain requirements in selecting data integration methods, proposing the classification of integration methods into database data models and semantic integration. This addresses the challenges of integrating geospatial big data from various sources, including the heterogeneity of data sources and the need to consider institutional, social, legal, and policy specifications. Additionally, it covers the classification of data integration into six categories and the use of linked data to revolutionize conventional approaches to data access, discovery, and integration.

B. DBMS for geospatial data

Geospatial data management within Database Management Systems (DBMS) is pivotal for efficiently handling spatial information encompassing maps, satellite images, and GPS coordinates. PostgreSQL with the PostGIS Extension is a notable choice, blending the robustness of PostgreSQL with the spatial capabilities of PostGIS. This amalgamation facilitates spatial indexing, supports diverse spatial data types, and furnishes spatial functions, rendering it apt for an array of geospatial applications [5].

Oracle Spatial and Graph, an integral component of the Oracle Database, offers advanced spatial data management features including spatial data types, indexing, and an extensive array of spatial functions. Its seamless integration with Oracle's database features enhances its efficiency in handling geospatial data. Similarly, Microsoft SQL Server, equipped with spatial extensions, provides effective geospatial data management by supporting spatial data types, indexing, and functions. Integration with Microsoft's ecosystem further bolsters its scalability, particularly in cloud-based geospatial solutions.

MongoDB distinguishes itself as a NoSQL database with native support for geospatial data through GeoJSON format. This flexibility enables MongoDB to store and query spatial data efficiently, especially beneficial for applications with substantial data volumes. Furthermore, GeoServer serves as a robust platform for sharing and editing geospatial data over the web. While not a conventional DBMS, its compatibility with various data sources and support for OGC protocols make it indispensable for web-based GIS applications, offering advanced cartographic capabilities.

The selection of the most suitable DBMS depends on factors like application requirements, scalability, budget, and existing technology infrastructure.

C. SQL versus NoSQL

The paper [2] evaluates the efficiency of SQL and NoSQL geospatial databases in a real application context, specifically using Microsoft Azure Cloud Service to deploy the application. The databases considered in the evaluation are Azure SQL Database and Azure DocumentDB. The paper also mentions the use of GeoServer and GeoTools, and provides a qualitative comparison of the functionalities of four state-of-the-art geospatial databases.

The application allows both authorities and citizens to generate a report at a specific location (point) and to retrieve all reports available on a given geographical region. The performed experiments allow drawing some general empirical conclusions about the two geospatial DBaaS-based databases benchmarked in this experiments. Both databases are characterized by good average response times and hence they can be used in the back-end of geospatial applications that select objects based on their location or submit geolocated objects.

The Figure 7 completely compares the SQL and NoSQL DBMS. The main difference that is highlighted by the information reported in Figure 7 is that relational databases (PostGIS and Azure SQL) implement more geospatial functionalities than NoSQL ones (MongoDB and DocumentDB).Moreover, they are also more tightly integrated and supported by the GeoServer software. In fact, the plug-in that is used to connect GeoServer with MongoDB and DocumentDB is in the unsupported branch of the current version of GeoTools [24]. The interoperability with the GeoServer, or with a similar software, is indispensable for geospatial applications that need to visualized maps and geolocated objects, especially when mobile devices are involved. Hence, in terms of functionalities, spatial relational databases are preferable (they are a more mature technology) because they allow performing complex geospatial queries and are wellintegrated with geographic information systems (GIS) as well as with map server software such as GeoServer. In the Section V, we will perform a set of experiment to evaluate also the efficiency of the selected databases in a real application scenario.

Below are the definistion and specificaton of some dbms that are well design for storing geospatial datas.

1) PostGIS: PostGIS is an extension for PostgreSQL, a powerful open-source relational database system. It enhances PostgreSQL with robust support for geospatial data, allowing it to store, query, and analyze spatial information efficiently. With PostGIS, PostgreSQL becomes a backend spatial database for Geographic Information Systems (GIS). It adds support for geographic objects, spatial indexing, and a wide range of spatial functions, making it suitable for diverse geospatial applications. PostGIS enables users to store spatial data types such as points, lines, and polygons, and perform spatial queries to extract relevant information from the data. Its integration with PostgreSQL's relational capabilities provides a comprehensive platform for managing both spatial and non-spatial data within the same database system.

2) Azure SQL Database: Azure SQL Database is a fully managed relational database service provided by Microsoft Azure. While it may not have native support for geospatial data like PostGIS, it can still handle such data effectively through the use of spatial extensions and functions. Microsoft SQL Server, which forms the basis of Azure SQL Database, includes spatial extensions that enable it to manage geospatial data efficiently. These extensions provide support for spatial data types, spatial indexing, and spatial functions, similar to those found in PostGIS. Azure SQL Database's integration with other Azure services, such as Azure Maps and Azure Spatial Anchors, further enhances its capabilities for geospatial data management and analysis, allowing users to build scalable

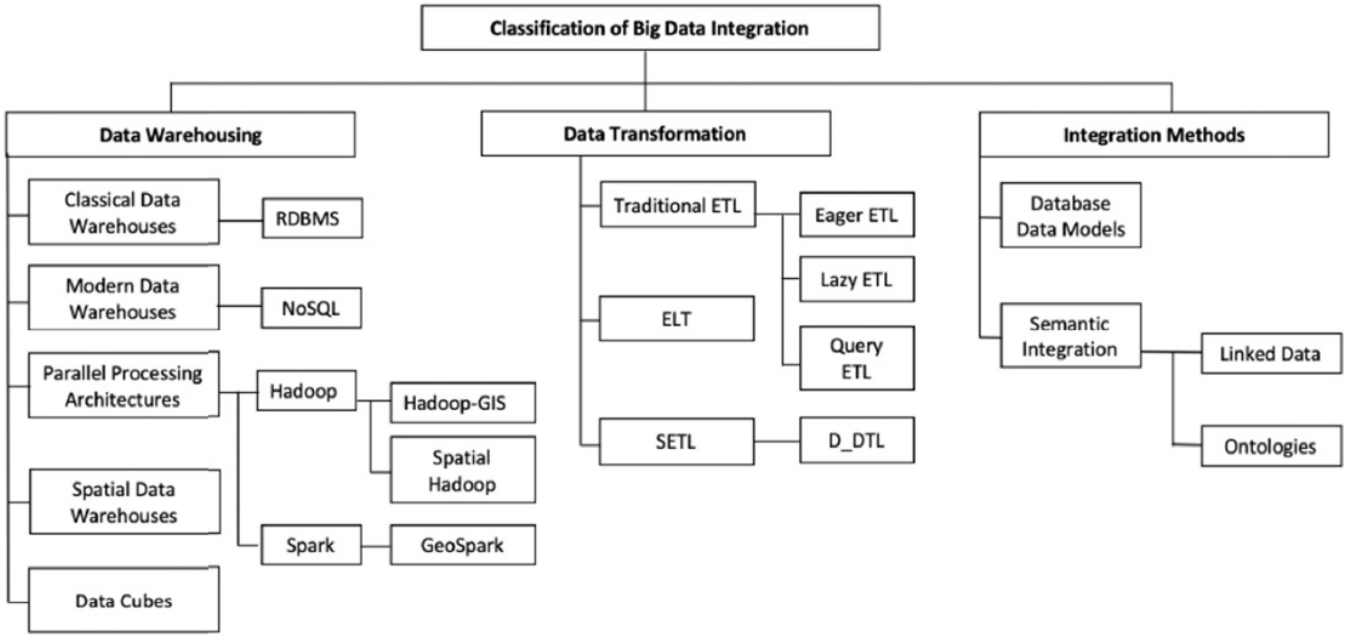

Fig. 6. Classification of big data integration.

Database	Supported Geometry objects	Main supported geometry functions	Supported Spatial indexes	Compatibility with GeoServer	DaaS	Horizontal scalability
PostGIS	Point, LineString, Polygon, MultiPoint, MultiLineString, MultiPolygon, GeometryCollection	PostGIS supports the Open Geospatial Consortium (OGC) methods on geometry instances	B-Tree index R-Tree index, GiST index	Yes	No	No
Azure SQL Database	Point LineString, Polygon, MultiPoint, MultiLineString, MultiPolygon, GeometryCollection	Azure SQL Database supports the Open Geospatial Consortium (OGC) methods on geometry instances	2d plane index, B-trees	Yes	Yes (Microsoft Azure cloud computing platform)	No
MongoDB	Point, LineString, Polygon, MultiPoint, MultiLineString, MultiPolygon, GeometryCollection,	Inclusion, Intersection, Distance/Proximity	2dsphere index, 2d index	Yes (based on the unsupported external MongoDB plug-in included in GeoTools)	Yes (MongoDB Atlas cloud service)	Yes (sharding)
DocumentDB	Point, LineString, Polygon, MultiPoint, MultiLineString, MultiPolygon, GeometryCollection	Inclusion, Distance/Proximity	2d plane index, quadtree	Yes (based on the unsupported external MongoDB plug-in included in GeoTools)	Yes (Microsoft Azure cloud computing platform)	Yes (sharding)

Fig. 7. QUALITATIVE COMPARISON OF THE FUNCTIONALITIES OF FOUR STATE-OF-THE-ART GEOSPATIAL DATABASES.

and cloud-based geospatial solutions.

3) MongoDB: MongoDB is a popular NoSQL database known for its flexibility and scalability. While it is not traditionally associated with geospatial data management, it offers support for such data through its GeoJSON format. GeoJSON allows MongoDB to store and query spatial data, including points, lines, and polygons, making it suitable for certain geospatial applications. MongoDB's flexible schema and horizontal scalability can be advantageous for handling large volumes of spatial data efficiently. MongoDB's spatial capabilities may not be as extensive as those of specialized geospatial databases like PostGIS, but it provides an alternative for applications that require a NoSQL approach to geospatial data management.

4) DocumentDB (Azure Cosmos DB): DocumentDB, now known as Azure Cosmos DB, is a globally distributed, multi-model database service provided by Azure. It supports various data models, including documents, key-value, graph, and column-family, making it versatile for diverse data management needs. While DocumentDB may not have native support for geospatial data, it can still handle such data by storing spatial information within documents or as part of complex data structures. Users can leverage DocumentDB's indexing and querying capabilities to work with geospatial data effectively. Azure Cosmos DB's global distribution and elastic scalability make it suitable for applications requiring geospatial data management across distributed environments or at scale.

Each of these database systems offers unique features and capabilities for managing geospatial data. Whether through native support like PostGIS, spatial extensions like Azure SQL Database, GeoJSON support like MongoDB, or flexible data modeling like DocumentDB, these databases can effectively handle spatial information to support a wide range of geospatial applications.

The paper [2] focuses on the evaluation of geospatial databases in the context of emergency management applications, providing insights into the performance and efficiency of SQL and NoSQL databases, as well as the use of specific web services and tools.

D. optimized interface programming language using python

The paper presents a Python-based framework called "ibmdbpy-spatial" for performing in-database geospatial analytics. It allows executing spatial queries and operations directly within a spatial database like IBM DB2, without having to load the entire dataset into memory. The method wraps spatial functions and queries into easy-to-use Python functions and methods that internally translate to SQL/OGC standard spatial queries. The spatial data is represented as a special "IdaGeoDataFrame" object, which extends the pandas DataFrame with spatial attributes and geometric operations. Users can write spatial analysis code in Python, which gets translated to database queries at runtime using SQL-pushdowns. The results are then fetched back as DataFrames. The framework leverages the performance benefits of in-database processing like columnar storage, indexing and parallel processing. It can handle different spatial data types like points, lines, polygons and supports common spatial operations like within(), intersects(), area(), length() etc. The framework can be used interactively within Jupyter notebooks for exploratory spatial data analysis and visualization using Python libraries. They demonstrate its applicability through a case study of analyzing crime hotspots in New York City using the in-database spatial functions. The Figure 8 shows the topological framework that underlies spatial operations for ibmdbpy-spatial.

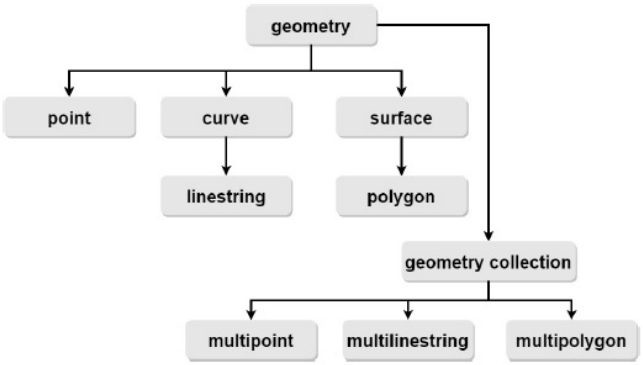

Fig. 8. Topological framework of spatial data used in ibmdbpy-spatial

E. GeoViz

GeoVisual analytics, abbreviated as GeoViz, is an emerging field that combines advanced data processing, visualization, and human-computer interaction techniques to enable effective analysis of large-scale geospatial datasets. It leverages the cognitive abilities of humans to comprehend complex spatial patterns and relationships represented through interactive visual interfaces [15].

The core process of GeoViz involves two key phases:

1) Spatial Data Preparation Phase: :

- Loading large geospatial datasets from diverse sources like databases (e.g., shapefiles, PostGIS) or distributed file systems (HDFS)
- Performing data processing operations like spatial filtering (range queries), spatial joins, geometric operations on the loaded datasets
- The output is a set of spatial objects (points, polygons, lines) relevant for visualization

2) Map Visualization Phase: :

- Applying visualization techniques like heatmaps, choropleth maps, etc. on the spatial objects
- Typically done by pixelizing the objects, aggregating overlapping pixels, color-coding, and rendering map tiles
- Enabling interactive exploration through pan, zoom, and visual overlays of multiple datasets

GeoViz systems distribute these phases across parallel computing clusters to scale to massive datasets that traditional desktop tools cannot handle efficiently. The computational challenges include data partitioning, parallel rendering, and optimization across the phases.

The output of GeoViz is an interactive visual environment that spatially represents the underlying data, enabling analysts to perceive patterns, trends, anomalies easily. It facilitates analytical reasoning by synergistically combining human vision, domain expertise and computational power.

GeoViz has applications across diverse domains like transportation, urban planning, environmental monitoring, public health, business analytics and more, providing critical insights for decision-making from complex geospatial big data.

F. Keyword Search

The paper [4] proposes a methodology for efficient storage and management of big geospatial data using HBase and MapReduce. The key components of the methodology are:

1) Data Layer: :

- Converts shapefile data into unified WKT (Well-Known Text) or WKB (Well-Known Binary) formats for vector geometries, following the OGC (Open Geospatial Consortium) Simple Features specifications. This allows interoperability with other systems.

2) Storage Layer: :

- Uses HDFS (Hadoop Distributed File System) as the underlying storage system for the HBase-based storage model.
- Utilizes GeoHash, a hierarchical spatial data structure, for spatial indexing of vector data.
- Designs the structure of HBase tables to efficiently store and manage vector data: a. Column Families based on vector data types (Point, LineString, Polygon, etc.) b. Columns to store different attributes of each data type c. Row Keys generated by converting average coordinates to GeoHash strings
- Employs MapReduce for efficient batch insertion of data into HBase, using the BufferedMutator interface provided by HBase.

3) Access Layer: :

- Provides a user-oriented API to write and read vector data to/from the HBase-based storage model, wrapping the native HBase API.

The methodology aims to leverage the distributed, scalable, and high-performance capabilities of HBase and MapReduce to store and manage large volumes of geospatial vector data efficiently. The use of GeoHash indexing and the design of HBase table structure facilitate efficient storage and retrieval of data based on spatial characteristics.

IV. Discussion

The presented paper offers a comprehensive review of geospatial databases, delving into their evolution, capabilities, challenges, and practical applications across various domains. It systematically examines the transition of geospatial data management from traditional paper maps to digital formats, highlighting the pivotal role of geospatial databases in modern data-driven decision-making processes. By exploring different types of geospatial databases, including relational, spatial, and NoSQL solutions, the paper elucidates their features, advantages, and limitations, providing valuable insights for researchers and practitioners in the field.

One of the significant contributions of the paper lies in its exploration of the challenges posed by geospatial big data, including issues related to volume, velocity, variety, and veracity. As the volume and diversity of geospatial data continue to grow exponentially, traditional relational database management systems (RDBMS) face limitations in efficiently storing, managing, and analyzing such data. To address these challenges, we discusses modern approaches like NoSQL databases and distributed processing frameworks, which offer scalability, flexibility, and performance advantages for handling geospatial big data effectively.

Moreover, the paper examines emerging trends and future directions in geospatial database research, such as geospatial video applications, in-database geospatial analytics, and benchmarking frameworks for evaluating the performance of geospatial NoSQL databases. These trends reflect the ongoing efforts to develop innovative solutions that can address the evolving needs of geospatial data management and analysis in diverse application domains.

The paper also provides a comprehensive overview of related work in the field, encompassing various studies on geospatial big data management, security concerns, spatial knowledge discovery, database comparison for geospatial applications, and benchmarking frameworks. By synthesizing insights from these studies, the paper contributes to a deeper understanding of the state-of-the-art techniques and methodologies in geospatial database research.

V. Conclusions

In conclusion, the paper offers a comprehensive review of geospatial databases, encompassing their evolution, capabilities, challenges, and applications. It highlights the importance of geospatial databases in managing and analyzing spatially-referenced data across diverse domains, ranging from urban planning to environmental management. By exploring different types of geospatial databases and addressing challenges related to geospatial big data, the paper provides valuable insights for researchers, practitioners, and decision-makers involved in spatial data management and analysis.

Furthermore, the paper identifies emerging trends and future research directions in geospatial database research, emphasizing the need for scalable, flexible, and efficient solutions to address the evolving demands of geospatial data management. By synthesizing insights from related work, the paper contributes to advancing the state-of-the-art in geospatial database research and facilitating the development of innovative solutions for managing and analyzing geospatial data effectively.

Overall, the paper serves as a comprehensive resource for understanding the evolution, capabilities, challenges, and applications of geospatial databases, and it lays the foundation for future research in this important and rapidly evolving field.

References

[1] Ando, K., Shimizu, K., Ito, N., Iwata, K., Uchitane, T., Mukai, N., Jiang, Y.: Design of geospatial database for analysis of road traffic accidents in aichi prefecture. In: 2022 Joint 12th International Conference on Soft

Computing and Intelligent Systems and 23rd International Symposium on Advanced Intelligent Systems (SCIS&ISIS). pp. 1–7. IEEE (2022)
[2] Baralis, E., Dalla Valle, A., Garza, P., Rossi, C., Scullino, F.: Sql versus nosql databases for geospatial applications. In: 2017 IEEE International Conference on Big Data (Big Data). pp. 3388–3397. IEEE (2017)
[3] Deng, H., Yue, P., Yu, D., Cao, Z., Liu, R., Wang, K.: A geospatial data and model hub for online geospatial analysis. In: 2023 11th International Conference on Agro-Geoinformatics (Agro-Geoinformatics). pp. 1–6. IEEE (2023)
[4] Gao, F., Yue, P., Wu, Z., Zhang, M.: Geospatial data storage based on hbase and mapreduce. In: 2017 6th International Conference on Agro-Geoinformatics. pp. 1–4. IEEE (2017)
[5] Ge, Y., Lin, V., Daum, M., Haynes, B., Cheung, A., Balazinska, M.: Demonstration of apperception: a database management system for geospatial video data. Proceedings of the VLDB Endowment **14**(12), 2767–2770 (2021)
[6] Kim, S., Kanwar, Y.S.: Geoycsb: A benchmark framework for the performance and scalability evaluation of nosql databases for geospatial workloads. In: 2019 IEEE International Conference on Big Data (Big Data). pp. 3666–3675. IEEE (2019)
[7] Loukili, Y., El Aissi, M.E.M., Benjelloun, S., Lakhrissi, Y., Chougrad, H., Ali, S.E.B., El Boushaki, A.: Analysis and exploitation of geospatial big data: State of art. In: 2020 IEEE 14th International Conference on Application of Information and Communication Technologies (AICT). pp. 1–6. IEEE (2020)
[8] Rajpoot, M.S., Patel, P.: A comparative study on various aspects of security of geospatial data. In: 2014 Fourth International Conference on Communication Systems and Network Technologies. pp. 708–712. IEEE (2014)
[9] Roy, A., Fouché, E., Morales, R.R., Möhler, G.: In-database geospatial analytics using python. In: Proceedings of the 2nd ACM SIGSPATIAL International Workshop on Advances on Resilient and Intelligent Cities. pp. 17–24 (2019)
[10] Scott, G.J., Backus, K., Anderson, D.T.: A multilevel parallel and scalable single-host gpu cluster framework for large-scale geospatial data processing. In: 2014 IEEE Geoscience and Remote Sensing Symposium. pp. 2475–2478. IEEE (2014)
[11] Shangguan, B., Yue, P., Cao, Z., Wang, B.: Towards a geospatial big data platform for geospatial information services. In: 2019 8th International Conference on Agro-Geoinformatics (Agro-Geoinformatics). pp. 1–4. IEEE (2019)
[12] Shea, G.Y., Cao, J.: Geo-planar indexing (gpi)–
[13] Wang, Z., Zhang, H., Sun, J.: Framework and case study on geospatial knowledge discovery based on spatial knowledge management. In: 2011 Eighth International Conference on Fuzzy Systems and Knowledge Discovery (FSKD). vol. 2, pp. 1280–1284. IEEE (2011)
[14] Xavier, E.M., Ariza-López, F.J., Urena-Camara, M.A.: A survey of measures and methods for matching geospatial vector datasets. ACM Computing Surveys (CSUR) **49**(2), 1–34 (2016)
[15] Yu, J.: Src: geospatial visual analytics belongs to database systems: the babylon approach. SIGSPATIAL Special **9**(3), 2–3 (2018)